

Volumetric Ray Tracing

1062¹¹OTH Regensburg, Germany

Abstract

—START OF REVIEW—

1. Introduction

In Computer Graphics, objects usually are represented as a set of geometric primitives (e.g. triangles), displaying the surface of the object in question. However, this approach is not always suitable. For example, if the original data representation of the object is volumetric data (which might be produced by medical 3d scans [Lev88]), the traditional rendering technique would necessitate the creation of an intermediate surface representation that might introduce unwanted artifacts. Another issue arises if the object has no well-defined surfaces to which geometric primitives could be fitted, such as a cloud or fog [KVH84]. In such cases, volumetric ray tracing might be used, a technique in which rays are cast through a volume which contains information about its optical properties (e.g. color and opacity), sampled at various points within the volume, accumulated and projected on a 2d image (see figure 1) [Lev88]. A description of this algorithm is presented in this paper, alongside several strategies for improving computational speed, such as (probabilistic) early termination of a ray, and sampling at lower resolutions within the volume.

2. Mathematical Theory Behind Volumetric Ray Tracing

Volumetric ray tracing follows the same basic principle as classical ray tracing in which an image is rendered by spanning a pixel plane in front of the camera and casting one or multiple rays through each pixel. That means for a point \vec{x} on the plane we cast a ray in direction $-\omega$ and calculate the amount of light \vec{x} receives from direction ω , called $L(\vec{x}, \omega)$. To do this, we calculate \vec{y} , the closest intersection point of the ray with a piece of geometry. At \vec{y} , we calculate the amount of light transported from \vec{y} in direction ω (either through reflection or emission), called $L_e(\vec{y}, \omega)$ (The exact method for calculating L_e is a topic of classical raytracing and will not be further evaluated in this paper. In the following, we assume L_e to be a known quantity). Thus, we get the equation

$$L(\vec{x}, \omega) = L_e(\vec{y}, \omega) \quad (1)$$

However, this equation assumes that the light has no interactions between \vec{x} and \vec{y} , which is only true if the ray travels through a vacuum. If the light travels through a medium that interacts with it (called a participating medium), the interactions change the light along the ray. In most cases, these interactions are small enough to be ignored, but in certain scenarios (e.g. if there is fog, clouds or smoke present) they might have a noticeable effect. The types of interactions we need to consider are absorption, emission, in-scattering and out-scattering.

Absorption and out-scattering attenuate the light intensity along the ray, in-scattering and emission add to it. Firstly, let us consider only out-scattering and absorption, which occur due to tiny particles floating around in the volume (such as water droplets in clouds and fog, or dust particles in the air). Of course, these particles are far too numerous to be directly simulated, but their distribution in 3D space can be stochastically modeled (similar to how detailed surface structures can be modeled by micromirrors in classical ray tracing).

To do so, consider a close-up look at a ray section traveling through the volume [CITATION!]. This section can be assumed to be a cylinder with a base area E and a height δh , through which the ray travels from top to bottom. Within this cylinder, there exists a certain number of out-scattering and absorbing particles, defined as $n_s = sE\delta h$ and $n_a = aE\delta h$ respectively, where s and a are the densities of the particles. From the top-down view, the particles will occupy an area of $\frac{n_s A = sAE\delta h}{E}$ and $\frac{n_a A = aAE\delta h}{E}$, respectively, if we assume that the particles do not overlap each other, which is reasonable if the densities and the height do not become too large. This can be simplified to $sA\delta h$ and $aA\delta h$, which gives us the fractions of light which are stopped in the cylinder, either through absorption or out-scattering. (1)

3. Basic Volumetric Ray Tracing Algorithm

3.1. Terminology

To enable a clear understanding in this section we will shortly define key terminology used in this paper. A voxel is an infinitesimally small point in 3D space with an associated vector value, that




Figure 1. Ray Casting scheme

[App15]

Figure 1: Illustration of the ray casting process.

represents the color of that voxel. In this paper, these voxels are arranged in a regular cube lattice (although other arrangements are possible as well). A cell or area of influence of a voxel v is defined as the region of space which is closer to v than to any other voxel (meaning a cube in space with v in the center). A voxel cube is defined as a group of 8 voxels that form a cube in space. The neighborhood of a voxel v is defined as the set that includes v as well as its 26 neighbors in 3D space. Since different authors use different notations, this terminology might be used slightly different in other papers.

3.2. Underlying Volume Data

In this paper, we consider the topic of synthesizing a 2d image from a 3d volume. A volume can either be some continuous, ternary function (such as perlin or simplex noise [Per85]), which can be used to render volumetric clouds [Häg18]), or a discrete 3d array [Lev88]. Similar to a 2d image that consists of pixels that can be addressed by a 2d vector $\vec{x} \in \mathbb{N}^2$, a 3d array consists of voxels that are addressed by a 3d vector $\vec{x} \in \mathbb{N}^3$. [REFORMULATE?]The color of a voxel at position \vec{x} is called $c(\vec{x})$, the opacity $\alpha(\vec{x})$ and the complete vector $v(\vec{x})$. The opacity is a scalar value, the color may also be a scalar value (for grayscale volumes) or a 3d vector (for colorful volumes). To address the components of a vector \vec{v} , following notation is used:

$$\vec{v}_x, x \in sr, g, b, \alpha \quad (2)$$

In the following, all values are assumed to be normalized to between 0 and 1. To project this volume to a image, for each pixel in the image a ray is cast through the volume and sampled at multiple, evenly spaced points on the ray (see figure 1) [Lev90].

3.3. Volumetric Ray Tracing

Like in the classical ray tracing algorithm, volumetric ray tracing works by casting a ray (or multiple rays if multisampling is used) for each pixel in the image that is to be created.

These rays are described by the equation $\vec{o} + t \cdot \vec{d}$, where \vec{o} is the origin of the ray and \vec{d} the direction. On each ray evenly spaced sample points are placed whose position is described by

$$\vec{o} + n \cdot s \cdot \vec{d} \quad (3)$$

for the n -th sample point. s is a scale factor, determining how far

apart the sample points are. s should be roughly equal to the distance between voxels (this distance is assumed to be 1 in the volume model, but depending on the volumes world matrix this might be different in world coordinates), since a too great mismatch between sampling and voxel frequency would lead to aliasing. In the following, only sample points within the volume are considered. Those sample points then are sampled and composited together, resulting in a final color value for the image pixel.

3.4. Interpolation

The sample points on the ray are in general part of \mathbb{R}^3 . This is no problem if the volume is described by a continuous function, which is defined everywhere. However, if it is a discrete 3d array only defined for points in \mathbb{N}^3 , the value of the volume at the sample point must be interpolated. This interpolation is done over the 8 closest voxels to the sample point (usually with a trilinear interpolation). The opacity values can be interpolated regularly, but the color values must be weighted with the respective opacity values before interpolation [WMG98]. To see why this is necessary, consider figure 2, which presents a simplified 2d example of a completely opaque, white object behind completely transparent empty space. Naively interpolating in this volume would result in two sample points, one completely white and opaque, the other gray and semitransparent. This gray point is not present in the original data and is an unwanted artifact. Weighting the colors with their opacity prevents this from happening.




Figure 2: Simplified visualization of naive interpolation in 2 dimensions. The black dots represent the voxels, the red dots the sample points. Notice that the second sample point has a gray color, even though the original data only has white voxels and black, transparent voxels, which shouldn't affect the final color. Adapted from Wittenbrink et al. [WMG98]

3.5. Compositing Of Multiple Sample Points

To compose the various sampled points on the ray to a single pixel, the points one after the other are alpha blended together. This compositing can be done front to back [Sab88, UK88] (starting with the sample point closest to the camera, blending it with the second, then the third, and so on), or back to front [DCH88, Lev88] (vice versa). Usually, the back to front approach is chosen since it allows to optimize the computation [Lev90] (see below). In the following, $c(i)$ is the color of the i -th sample point, and $\alpha(i)$ the opacity. The

accumulated opacity

$$\beta(i) = 1 - \prod_{j=1}^i (1 - \alpha(j)) \quad (4)$$

is the fraction of light absorbed between the first and i -th sample point. The final color c_f , which is displayed in the projected image, is then

$$c_f = \sum_{i=1}^n ((1 - \beta(i)) * c(i)) \quad (5)$$

if the ray has n sample points [DH92].

4. Optimization Strategies

4.1. Adaptive Termination Of Ray Tracing

When looking at the compositing process described above, it becomes apparent that the contribution of a sample point to the final result decreases, the more other opaque sample points lie before it [Lev90]. For an extreme example, consider figure 3. Here, all sample points behind the fully opaque black wall contribute nothing to the final value and can therefore be ignored. In other words, going back to equation (2), the ray casting can be terminated after the accumulated opacity reaches 1, without changing the image quality. If some errors in the image are acceptable, the threshold value for terminating the ray tracing might be chosen to be somewhat lower.




Figure 3: The black, opaque object blocks parts of the volumes behind it. The samples behind the object (from the cameras perspective) are not visible and therefore irrelevant.

4.2. Ray Termination With Russian Roulette

The above described termination after a certain threshold is reached creates a bias in the synthesized image [AK90]. One way to avoid this is to use Russian Roulette to decide when to terminate a ray [DH92]. Unlike the above described approach, Russian Roulette terminates a ray only with a certain likelihood once the accumulated opacity reaches the threshold. The weight of the surviving rays is increased proportionally to compensate for the terminated rays.

4.3. Pyramid Data Structures

The following optimizations require a data structure called a pyramid. Pyramids are created from 3d arrays, if the volume to be rendered is defined as a function an intermediate array representation must be created. Analogous to a mip-map which is a set of 2d arrays with decreasing resolution, a pyramid is a set of 3d arrays with decreasing resolution. The different volumes of which the pyramid consists are called the levels of the pyramid. The lowest level (called level 0) has the highest resolution, each successive level has half the resolution than the preceding one. To ensure that each level covers the same region in 3d space, the distance between the voxels are doubled each level.

–FIGURE! In the following, we will make use several different kinds of pyramids, namely average, maximum and range pyramids.

4.4. Average Pyramid

The lowest level of the average pyramid is equal to the original 3d array, padded with zeros so that its size is a power of 2 in all dimensions. The next level is created by aggregating the average of a cube of 8 voxels in the lower level to one voxel in the higher level. The voxel in the higher level is called the parent, the 8 lower voxels are called the children. This process is repeated until the last level, consisting of only one voxel, is reached. Similarly to the original volume, the pyramid can be sampled by interpolating the 8 closest voxels to the sample point. Consider however, that the primary purpose of the average pyramid is to enable an approximation of the original data that is computationally faster than to render the original data. The time complexity of rendering 3D data however does not depend on the number of voxels, but rather the number of samples along the ray (at first glance, this seems to imply that the average pyramid is superfluous and that the same could be accomplished by taking fewer sampling points in the original volume. This however would lead to a sample frequency significantly lower than the voxel frequency and thus to aliasing). If we therefore wish to cut the rendering time in half for each level we move up the pyramid, the amount of sampling points must be halved and the distance between the sampling points must be doubled. This means that one sample in level $n + 1$ should approximate two samples at level n . To achieve this, the sample at level $n + 1$ must be blended with itself, since in level n two samples are blended as well.

4.5. Maximum And Minimum Pyramids

A maximum pyramid works similar to an average pyramid, but the construction of level 0 is somewhat different. A voxel in level 0 at position \vec{x} is created by taking the maximum of the 27 voxels in the original dataset, that form a 3 by 3 cube around \vec{x} . Once the lowest level is created the succeeding levels are constructed like in the average pyramid, only that the maximum is used as the aggregation method, and not the maximum. Sampling a maximum pyramid is different from sampling an average pyramid. Since the maximum pyramid stores maxima, interpolating between different maxima would lead to wrong results. Instead, nearest neighbor interpolation is used (unlike in the average pyramid, the sample does not need to be blended with itself). The minimum pyramid works analogously. In the following, the sample taken at location \vec{x} in the

maximum pyramid at level n is called $\text{pyr}_{\max}(\vec{x}, n)$, and the sample in the minimum pyramid is called $\text{pyr}_{\min}(\vec{x}, n)$.

4.6. Range Pyramid

The range pyramid is calculated by taking the difference between the maximum and minimum pyramid. Both the maximum and minimum pyramid store vectors (2D or 4D, depending on whether the underlying data is grayscale or colorful). To express the difference in a scalar, the difference of the components is summed up and normalized to the range of 0 and 1:

$$\text{pyr}_{\text{range}}(\vec{x}, n) = \frac{4 \cdot ((\text{pyr}_{\max}(\vec{x}, n)_r - \text{pyr}_{\min}(\vec{x}, n)_r) + (\text{pyr}_{\max}(\vec{x}, n)_g - \text{pyr}_{\min}(\vec{x}, n)_g) + (\text{pyr}_{\max}(\vec{x}, n)_b - \text{pyr}_{\min}(\vec{x}, n)_b) + (\text{pyr}_{\max}(\vec{x}, n)_a - \text{pyr}_{\min}(\vec{x}, n)_a))}{4} \quad (6)$$

The range pyramid is sampled like the maximum and minimum pyramids and provides a measure of the homogeneity of the original data. If a voxel in the range pyramid has the value 0, this means that the region of space covered by this pixel is completely homogenous in the original data, if the voxel is 1, this means at least 2 voxels in the original data are completely different.




Figure 4: Simplified 2 dimensional example of 2 levels of a pyramid. Each pixel in level $n + 1$ contains 4 pixel of the lower level. In 3 dimensions, 8 voxels of the lower level are contained in one voxel of the upper level.

4.7. Fixed Step Multiresolution

This optimization makes use of the above described average pyramid, and works by casting the rays through a level of the pyramid, instead of the original data. As previously mentioned, when casting a ray through level n , the distance between the sampling points is 2^n times larger than the normal distance. The level through which the rays are cast can be freely chosen. Similar to mip-maps, it is also possible to select a real number as a level, in this case, the levels above and below the chosen level are rendered, and the two results are interpolated together. Since this would require rendering the pyramid at two levels, this is slower than just choosing a natural number as a level and recommendable only if a smooth transition between two levels is needed, like in gaze directed rendering as described by Levoy(CITE!!). The higher the level the faster the algorithm works, but the lower the resolution becomes.

4.8. Presence Acceleration

This algorithm uses an average and a maximum pyramid. The maximum pyramid is used to quickly find regions with low opacity,

where the average pyramid is used to sample at a lower resolution. The idea behind this algorithm is that low opacity regions do not contribute much to the final result, and can be therefore be sampled more sparsely. The algorithm works by casting each ray through the maximum pyramid at the highest level. For each cell that is intersected, it checks if the opacity value of that cell is larger than some user provided threshold $k \in [0, 1]$. If so, that means that the average opacity in that region is not small enough, and we move down one level in the pyramid in the hope of finding such a region at a lower resolution. However, if the opacity is smaller the average pyramid is used to approximate that cell. [MORE DETAILS?]
Once a cell in the maximum pyramid has been checked, the algorithm advances to the next cell and checks if the new cell has the same parent as the old one. If not, we move up one level in the pyramid. This is done to ensure that the algorithm always advances with the largest possible step size. After that, the same procedure is repeated until the ray has moved through the entire pyramid. Further improvements to this algorithm can be made by considering two observations: Firstly, it is very rare for the cell in the highest level of the maximum pyramid to have a α lower than k (this would mean that the entire volume is almost completely transparent). It is therefore preferable not to start at the highest level but somewhat lower. XXX suggests in [CITE] to start two levels lower. Secondly, finding the cells that the ray intersects is a computationally relatively expensive operation that might not amortize itself at lower levels. Therefore, it might be more efficient to start sampling at the lowest level of the average pyramid before the lowest level of the maximum pyramid is reached. XXX suggests in [] to sample level 0 of the average pyramid once level 2 of the maximum pyramid is reached. [REFORMULATE?]

4.9. Homogeneity Acceleration

This algorithm works conceptually the same as presence acceleration. The difference is that homogeneity acceleration takes fewer samples in regions with high homogeneity and not in regions of low opacity. The idea behind the algorithm is that in a region where all voxels are very similar to each other, accumulating x different samples and accumulating the average of the region x times yields a very similar result. The algorithm works in the same way as presence acceleration, only the maximum pyramid is replaced with the range pyramid.

4.10. β Acceleration

References

- [AK90] ARVO J., KIRK D.: Particle transport and image synthesis. *SIGGRAPH Comput. Graph.* 24, 4 (sep 1990), 63–66. URL: <https://doi.org/10.1145/97880.97886>, doi:10.1145/97880.97886. 3
- [App15] APPA S.: Ray casting for 3d rendering – a review. 2
- [DCH88] DREBIN R. A., CARPENTER L., HANRAHAN P.: Volume rendering. *SIGGRAPH Comput. Graph.* 22, 4 (jun 1988), 65–74. URL: <https://doi.org/10.1145/378456.378484>, doi:10.1145/378456.378484. 2
- [DH92] DANSKIN J., HANRAHAN P.: Fast algorithms for volume ray tracing. In *Proceedings of the 1992 Workshop on Volume Visualization*

- (New York, NY, USA, 1992), VVS '92, Association for Computing Machinery, p. 91–98. URL: <https://doi.org/10.1145/147130.147155>, doi:10.1145/147130.147155. 3
- [Häg18] HÄGGSTRÖM F.: Real-time rendering of volumetric clouds, 2018. 2
- [KVH84] KAJIYA J. T., VON HERZEN B. P.: Ray tracing volume densities. *SIGGRAPH Comput. Graph.* 18, 3 (jan 1984), 165–174. URL: <https://doi.org/10.1145/964965.808594>, doi:10.1145/964965.808594. 1
- [Lev88] LEVOY M.: Display of surfaces from volume data. *IEEE Computer Graphics and Applications* 8, 3 (1988), 29–37. doi:10.1109/38.511. 1, 2
- [Lev90] LEVOY M.: Efficient ray tracing of volume data. *ACM Trans. Graph.* 9, 3 (jul 1990), 245–261. URL: <https://doi.org/10.1145/78964.78965>, doi:10.1145/78964.78965. 2, 3
- [Per85] PERLIN K.: An image synthesizer. *SIGGRAPH Comput. Graph.* 19, 3 (jul 1985), 287–296. URL: <https://doi.org/10.1145/325165.325247>, doi:10.1145/325165.325247. 2
- [Sab88] SABELLA P.: A rendering algorithm for visualizing 3d scalar fields. *Proceedings of the 15th annual conference on Computer graphics and interactive techniques* (1988). 2
- [UK88] UPSON C., KEELER M.: V-buffer: visible volume rendering. *Proceedings of the 15th annual conference on Computer graphics and interactive techniques* (1988).
- [WMG98] WITTENBRINK C., MALZBENDER T., GOSS M.: Opacity-weighted color interpolation for volume sampling. In *IEEE Symposium on Volume Visualization (Cat. No.989EX300)* (1998), pp. 135–142. doi:10.1109/SVV.1998.729595. 2